

WORKING DOCUMENT

Your LinkedIn Launch Pack

The Presence Kit filled in for the founder/CEO track: positioning options, profile copy, and the first nine posts.

Contents

1. How to use this pack

2. Positioning draft

3. LinkedIn profile copy

4. The first nine posts

Figures for growth, timelines and prices in this document are planning estimates, not measured research. Test them against your own results and adjust.

How to use this pack

Subject: the repo owner, founder/CEO track, LinkedIn, 3–5 hrs/week.

FILE	WHAT IT IS	STATUS
positioning-draft.md	Three positioning options, with a recommendation	Needs your decision
linkedin-copy.md	Paste-ready headline, About, Featured, settings	Draft — blocked on positioning choice + your background
first-9-posts.md	The nine launch posts, hooks written	Draft — needs your real examples in the slots

Read this before using any of it

I wrote the draft from what's visible in this repository, not from facts you gave me. **Anything marked [CONFIRM] is my inference and may be wrong.** The background paragraph in the About section is deliberately left blank rather than invented — a fabricated credential on a public profile is the one mistake you cannot walk back.

The five things I need from you

Answer these and I'll finish all three files in one pass:

1. **Your name**, exactly as it should appear, and your **current role and company**.
2. **Which positioning option** — A (applied-AI builder, my recommendation), B (African-market technologist), or C (builder in public).
3. **Your background in 3–4 lines**: what you've built, where you've worked, anything a stranger would recognise. If you're early and don't have credentials yet, say so — "early and shipping" is a workable position; vagueness is not.
4. **The one action** a follower should take: contact you for work, join a newsletter, or join a waitlist.
5. **Two real stories** for the posts — ideally something you refused to ship (post 4) and something you got wrong (post 8). These two carry the whole set.

Then the order of operations

1. Confirm the five above → I finish the copy
2. You take the photo (20 min, window light, plain wall — see `../IMPLEMENTATION.md`)
3. Profile goes live
4. Nine posts banked, three scheduled
5. Build the list of 20 target accounts
6. Announce, then start the weekly rhythm

Realistically: profile live within a week, first post the week after. Days 30 / 60 / 90 targets are in `../playbooks/ceo.md`.

Positioning draft

This is a draft I wrote from what's visible in this repo, not from facts you gave me. Everything marked **[CONFIRM]** is my inference and may be wrong. Correct it and I'll rewrite the profile and posts to match. Nothing here goes live until you've checked it.

What I'm inferring about you

From the work in this repository:

SIGNAL	WHAT I READ FROM IT
A LAMP diagnostic assistant with hard clinical safety rules	You build real tools in high-stakes domains, and you care about them being safe, not just impressive
Adire indigo textile patterns, Yoruba technique named correctly	Heritage and craft matter to you, and you handle them with specificity rather than decoration
A social-media system packaged for resale, with honest milestone numbers	You think in systems and products, not one-offs
Working across diagnostics, culture, and AI in one week	You're a builder who ships across domains fast

- **[CONFIRM] Name**, as you want it to appear: _____
- **[CONFIRM] Current role and company**: _____
- **[CONFIRM] Where you're based**: _____

Three positioning options

Pick one. They're mutually exclusive in practice — trying to be all three means being none of them.

Option A — The applied-AI builder (my recommendation)

I build AI tools for work where being wrong has a cost — diagnostics, health, and frontline operations — and I write about what it actually takes to make them safe enough to use.

Why I'd pick this one: it's the most defensible. Nearly everyone posting about AI right now is posting takes. Almost nobody is posting *working systems with safety constraints they can show*. The LAMP assistant is proof you can point at, and that's rare enough to be a position.

Audience: health-tech founders, NGO and public-health programme leads, funders, engineers in regulated domains.

What they should do: contact you to build or advise.

Option B — The African-market technologist

I build technology for African markets — diagnostics, commerce, and culture — and I write about what Western-built tools get wrong about this context.

Why it could work: strong point of view, underserved audience, and you clearly have real context. **Why I'd hold it:** it only lands if you have specific ground-truth stories (a deployment, a field trip, a failed assumption). If you have those, this beats Option A.

Option C — The builder-in-public generalist

I ship one working thing a week — AI tools, brands, systems — and I show the whole build.

Why I'd avoid it as the main line: it grows fast and converts badly. You get an audience of other builders, not buyers. Good as a *content format*, bad as a position. You can absolutely build in public *while* positioned as A.

If you pick Option A, here's the rest

One-liner:

I build AI tools for high-stakes work — diagnostics, health, frontline operations — where a confident wrong answer costs more than no answer.

Bio line (≤150 chars):

Building AI for work where being wrong has a cost. Diagnostics, health, frontline tools. I write about what makes them safe enough to actually use.

The belief that drives your content (this is your engine — every strong post comes back to it):

Most AI products are optimised to look confident. In the places where the stakes are real, the valuable behaviour is the opposite: knowing what you don't know, and saying so.

That belief is unusual, defensible, provable from your own work, and something a good number of people in AI would argue with. That's exactly what you want.

Content pillars

PILLAR	SHARE	WHAT GOES HERE
Building in high-stakes domains	35%	Safety constraints, failure modes, what you refused to ship and why
Point of view on AI	30%	Where the industry is wrong about reliability, evaluation, confidence
The work itself	20%	What you shipped, what broke, the numbers
The human	15%	Why this, where you're from, what you got wrong

Voice

YOU ARE	YOU ARE NOT
Specific, with real numbers and named constraints	Vague, "AI is transforming everything"

YOU ARE	YOU ARE NOT
Willing to say what didn't work	Only showing wins
Plain about limits	Hype-y or absolutist

No emoji bullets. No "Comment YES below." No rocket ships. Your whole position is credibility; engagement bait spends it.

The one action

Every profile and post points to: **[CONFIRM]** — pick one. - a contact line for build/advisory work - a newsletter - a waitlist for something you're building

Without this chosen, the whole thing generates attention and no outcome.

LinkedIn profile copy

Written for **Option A** (applied-AI builder). If you pick B or C, tell me and I'll rewrite. [...] means I need a fact from you.

Headline (220 char limit)

Pick one:

1 — plainest, my pick

I build AI tools for high-stakes work – diagnostics, health, frontline ops. Writing about what makes them safe enough to actually use.

2 — with the company

Founder, [Company] · AI tools for diagnostics and frontline health work · Writing about reliability in places where a wrong answer costs

3 — belief-forward

Building AI that knows what it doesn't know · Diagnostics & frontline health tools · [Company]

About section

Opening line is everything — it's all that shows before "...see more".

Most AI products are built to sound confident. I work on the ones where that's the dangerous setting.

I build tools for high-stakes work – diagnostics, health, frontline operations – where a confident wrong answer costs more than no answer at all. That changes what you optimise for. A diagnostic assistant I built refuses to report any sample result when the run's controls failed, no matter how clear the sample looks, because that's the most common way a field result goes wrong. Refusing is the feature.

[CONFIRM – 2-3 lines of your actual background: what you've built, where you've worked, anything a stranger would recognise or that proves you've done this before. If you're early and don't have credentials yet, say what you're building instead. Being early is fine; vagueness isn't.]

I write here about what it takes to make these systems trustworthy: the failure modes, the constraints worth hard-coding, and the gap between a demo and something a nurse or technician can rely on.

If you're building in a domain where being wrong has a cost, I'd like to hear about it – [your email or "DM me"].

Featured section (4 slots)

1. **The LAMP assistant** — link the repo or a short write-up. This is your strongest proof. Caption: *"A diagnostic assistant that refuses to answer when the controls failed."*
 2. **[CONFIRM]** Your best previous work or company site
 3. **[CONFIRM]** Any talk, interview, or press
 4. A contact or booking link
-

Experience entries

Rewrite each role as 2–3 lines of **outcomes**, not duties.

Bad: *"Responsible for leading development of diagnostic software."* Good: *"Built and shipped [tool] for [who]. [Number] users / [number] sites. Cut [thing] from X to Y."*

[CONFIRM] — **send me your roles and I'll write these.**

Settings checklist

- ☐ Profile set to public
- ☐ "Follow" as the primary button, not "Connect", so strangers can follow
- ☐ Custom URL: `linkedin.com/in/[yourname]`
- ☐ Banner: not a stock skyline. Either you speaking/working, or plain brand colour with your one-liner in text
- ☐ Contact info includes a real email
- ☐ 2FA on

The first nine posts

Bank all nine **before** you announce the profile. Schedule three; keep six in reserve so a busy week can't break the streak.

How to use these: the hooks are written and should stay close to as-is — the first two lines do all the work on LinkedIn. The bodies are scaffolded, and [. . .] marks where your real detail goes. **A post with my generic example in it is worse than no post.** If you can't fill a slot with something true, cut that post and use a reserve.

Format rules: short lines, blank line between each, no emoji bullets, links in the first comment.

1 · The belief (post this first)

Pillar: point of view

Most AI products are built to sound confident.
I work on the ones where that's the dangerous setting.

Body: - The domains you work in and why confidence is a liability there - One concrete example of a system refusing to answer, and why that was correct - [your own line on what you think the industry gets wrong]

Close: *What's a case in your work where "I don't know" was the right output?*

2 · The controls story

Pillar: building in high-stakes domains · your strongest proof

A diagnostic test can give you a perfectly clear positive result that means absolutely nothing.
Here's the rule we hard-coded because of it.

Body: - If the negative control shows amplification, the run is contaminated. Every tube in it is unreadable, however clean the sample looks - So the assistant checks controls first and returns INVALID for the whole run, refusing to report any sample result - The lesson: in high-stakes tools, the valuable behaviour is often the refusal, not the answer - [what this cost you to build, or what a reviewer said about it]

Close: *What's the equivalent "check the controls first" rule in your domain?*

3 · The demo gap

Pillar: point of view

The distance between an AI demo and something a nurse will actually use is not a engineering gap.
It's a trust gap, and it's mostly made of edge cases nobody filmed.

Body: three things that break between demo and deployment in your experience — [be specific: lighting, connectivity, who's holding the phone, what happens at 3am]

Close: one line on what you do differently now.

4 · What I refused to ship

Pillar: building in high-stakes domains

[The feature someone asked you for that you said no to]
I said no, and I'd say no again.

Body: what was asked, why it was reasonable to ask, the specific harm it would have enabled, what you built instead.
This is your most credible post type. It's also the hardest to fake, which is exactly why it works.

5 · Where I'm from

Pillar: the human

[Opening about your background / where you grew up / what you saw]

Body: the connection between that and what you build now. Specific scene, not a mission statement.
Optional and strong if it's true: the Adire indigo work — a craft where the pattern is made by what you *block out*, not what you add. Resist-dyeing as a way of thinking about constraints in a system. Only use this if it's genuinely yours; borrowed heritage reads instantly as borrowed.

6 · The uncomfortable opinion

Pillar: point of view · expect disagreement

Benchmark scores tell you almost nothing about whether a model is safe to put in front of a technician.
[Your sharper version of this]

Body: why the thing the industry measures isn't the thing that matters in your domain, what you'd measure instead, and the one case that convinced you.
Close with a real question, not a rhetorical one. Reply to every disagreement in the comments — that's where this post earns its value.

7 · The ship post

Pillar: the work

[What you shipped] . Here's what's in it and what's still wrong with it.

Body: what it does, one number, one honest limitation, what's next. The limitation is what makes people trust the rest.

8 · The mistake

Pillar: the human

[A real thing you got wrong – a wrong assumption, a feature nobody used, a deployment that failed]

It cost [time/money/trust] . Here's what I'd do differently.

Body: what you believed, what actually happened, what changed in how you work.

No redemption arc. Just the lesson.

9 · The open door

Pillar: the work · your one direct ask of the nine

I'm looking to work with [specific type of team] on [specific type of problem] .

Body: two sentences on what you bring, one on what a good fit looks like, one on what to do next.

One ask in nine posts. That ratio is what makes the ask land.

After these nine

Your ongoing rotation, 3 a week:

WEEK	MON	WED	THU
A	Building / high-stakes	Point of view	The work
B	Point of view	Building / high-stakes	The human

Repeat. Pull topics from what you actually did that week — the best posts are Tuesday's problem written down on Wednesday.